

# Kafka Cluster Topic Key Distribution Analyzer Tool

Efficient **Kafka key distribution** is fundamental to building scalable, high-performance event-driven systems. Kafka uses each record's key to determine which partition it belongs to—governing **data ordering**, **load balancing**, and **parallelism** across the cluster. When key distribution is uneven, some partitions become hot, processing far more traffic than others. These **hot partitions** lead to broker overload, consumer lag, and throttled throughput, undermining the scalability of your Kafka workloads.

This tool helps you **test**, **visualize**, and **validate** how record keys are distributed across topic partitions in your Kafka cluster. It generates records using configurable key patterns, publishes them to a target topic, and then consumes the data to analyze partition utilization and message distribution metrics.

By surfacing patterns of **data skew**, **low-key cardinality**, or **biased hashing**, the analyzer reveals whether your partitioning strategy is truly balanced. The results empower you to:

- Detect and diagnose **hot partitions** before they degrade performance.
- Experiment with **key-salting** or **hashing strategies** to improve balance.
- Optimize **consumer parallelism** and **broker load** for predictable throughput at scale.

Use this tool as a **proactive performance lens** on your Kafka topics—ensuring your cluster's data distribution is as efficient, scalable, and reliable as the workloads it powers.

## Table of Contents

- **1.0 To get started**
  - **1.1 Download the Tool**
  - **1.2 Configure the Tool**
    - **1.2.1 Create a Dedicated Service Account for the Analyzer Tool**
    - **1.2.2 Create the .env file**
    - **1.2.3 Using the AWS Secrets Manager (optional)**
  - **1.3 Run the Tool**
    - **1.3.1 Did you notice we prefix `uv run` to `streamlit run src/tool.py`?**
    - **1.3.2 Troubleshoot Connectivity Issues (if any)**
  - **1.4 The Results**
    - **1.4.1 Example of Hot Key Data Skew Simulation Results**
    - **1.4.2 Example of Normal Key Distribution Simulation Results**
    - **1.4.3 Example of Low Cardinality Key Distribution Simulation Results**
    - **1.4.4 Example of High Cardinality Key Distribution Simulation Results**
    - **1.4.5 Example of No Repetition Key Distribution Simulation Results**
- **2.0 How the Tool Works**
  - **2.1 End-to-End Flow**
- **3.0 Resources**

## 1.0 To get started

**Download** ---> **Configure** ---> **Run** ---> **Results**

### 1.1 Download the Tool

Clone the repo: `shell git clone https://github.com/j3-signalroom/kafka_cluster-topic-key_distribution_analyzer-tool.git`

Since this project was built using **uv**, please **install** it, and then run the following command to install all the project dependencies:

```
uv sync
```

### 1.2 Configure the Tool

Now, you need to set up the tool by creating a **.env** file in the root directory of your project. This file will store all the essential environment variables required for the tool to connect to your Confluent Cloud Platform and function correctly. Additionally, you can choose to use **AWS Secrets Manager** to manage your secrets.

#### 1.2.1 Create a Dedicated Service Account for the Analyzer Tool

The service account needs to have **OrganizationAdmin**, **EnvironmentAdmin** or **CloudClusterAdmin** role to provision Kafka cluster API keys and the **MetricsViewer** role to access the Metrics API for all clusters it has access to.

1. Use the **Confluent CLI (Command-Line Interface)** to create the service account:

**Note:** If you haven't already, install the **Confluent CLI** and log in to your Confluent Cloud account using `confluent login`. Moreover, the account you use to log in must have the **OrganizationAdmin** role to create the **Cloud API key in Step 5**.

```
confluent iam service-account create <SERVICE_ACCOUNT_NAME> --description "<DESCRIPTION>"
```

For instance, you run `confluent iam service-account create recommender-service-account --description "Service account for Recommender Tool"`, the output should resemble:

```
+-----+-----+
| ID      | sa-abcd123 |
| Name    | recommender-service-account |
| Description | Service account for |
|         | Recommender Tool   |
+-----+-----+
```

2. Make note of the service account ID in the output, which is in the form `sa-xxxxxxx`, which you will assign the [OrganizationAdmin](#), [EnvironmentAdmin](#) or [CloudClusterAdmin](#) role, and [MetricsViewer](#) role to in the next steps, and assign it to the `PRINCIPAL_ID` environment variable in the `.env` file.
3. Decide at what level you want to assign the [OrganizationAdmin](#), [EnvironmentAdmin](#) or [CloudClusterAdmin](#) role to the service account. The recommended approach is to assign the role at the organization level so that the service account can provision API keys for any Kafka cluster in the organization. If you want to restrict the service account to only be able to provision API keys for Kafka clusters in a specific environment, then assign the [EnvironmentAdmin](#) role at the environment level. If you want to restrict the service account to only be able to provision API keys for a specific Kafka cluster, then assign the [CloudClusterAdmin](#) role at the cluster level.

For example, to assign the [EnvironmentAdmin](#) role at the environment level:

```
confluent iam rbac role-binding create --role EnvironmentAdmin --principal User:<SERVICE_ACCOUNT_ID> --
environment <ENVIRONMENT_ID>
```

Or, to assign the [CloudClusterAdmin](#) role at the cluster level:

```
confluent iam rbac role-binding create --role CloudClusterAdmin --principal User:<SERVICE_ACCOUNT_ID> --cluster
<KAFKA_CLUSTER_ID>
```

For instance, you run `confluent iam rbac role-binding create --role EnvironmentAdmin --principal User:sa-abcd123 --environment env-123abc`, the output should resemble:

```
+-----+-----+
| ID      | rb-j3XQ8Y |
| Principal | User:sa-abcd123 |
| Role     | EnvironmentAdmin |
+-----+-----+
```

4. Assign the [MetricsViewer](#) role to the service account at the organization, environment, or cluster level, For example to assign the [MetricsViewer](#) role at the environment level:

```
confluent iam rbac role-binding create --role MetricsViewer --principal User:<SERVICE_ACCOUNT_ID> --environment
<ENVIRONMENT_ID>
```

For instance, you run `confluent iam rbac role-binding create --role MetricsViewer --principal User:sa-abcd123 --environment env-123abc`, the output should resemble:

```
+-----+-----+
| ID      | rb-1GgVMN |
| Principal | User:sa-abcd123 |
| Role     | MetricsViewer |
+-----+-----+
```

5. Create an API key for the service account:

```
confluent api-key create --resource cloud --service-account <SERVICE_ACCOUNT_ID> --description "<DESCRIPTION>"
```

For instance, you run `confluent api-key create --resource cloud --service-account sa-abcd123 --description "API Key for Recommender Tool"`, the output should resemble:

```
+-----+-----+
| API Key | 1WORLDABCDEF70AB |
+-----+-----+
```

| API Secret | cfltabCdeFg1hI+/2j34KLMnoprSTuvxy/Za+b5/6bcDe/7fGhIjkLMn0PQ8rT9U |

+-----+

6. Make note of the API key and secret in the output, which you will assign to the `confluent_cloud_api_key` and `confluent_cloud_api_secret` environment variables in the `.env` file. Alternatively, you can securely store and retrieve these credentials using AWS Secrets Manager.

1.2.2 Create the `.env` file

Create the `.env` file and add the following environment variables, filling them with your Confluent Cloud credentials and other required values:

```
# Environment variables credential for Confluent Cloud
CONFLUENT_CLOUD_CREDENTIAL={"confluent_cloud_api_key":"<YOUR_CONFLUENT_CLOUD_API_KEY>", "confluent_cloud_api_secret": "<YOUR_CONFLUENT_CLOUD_API_SECRET>"}

# Environment and Kafka cluster filters (comma-separated IDs)
# Example: ENVIRONMENT_FILTER="env-123,env-456"
# Example: KAFKA_CLUSTER_FILTER="lkc-123,lkc-456"
ENVIRONMENT_FILTER=<YOUR_ENVIRONMENT_FILTER, IF ANY>
KAFKA_CLUSTER_FILTER=<YOUR_KAFKA_CLUSTER_FILTER, IF ANY>

# Confluent Cloud principal ID (user or service account) for API key creation
# Example: PRINCIPAL_ID=u-abc123 or PRINCIPAL_ID=sa-xyz789
PRINCIPAL_ID=<YOUR_PRINCIPAL_ID>

# AWS Secrets Manager Secrets for Confluent Cloud and Kafka clusters
USE_AWS_SECRETS_MANAGER=<True|False>
CONFLUENT_CLOUD_API_SECRET_PATH={"region_name": "<YOUR_SECRET_AWS_REGION_NAME>", "secret_name": "
<YOUR_CONFLUENT_CLOUD_API_KEY_AWS_SECRETS>"}
```

The environment variables are defined as follows:

Environment Variable Name	Type	Description	Example	Default	Required
ENVIRONMENT_FILTER	Comma-separated String	A list of specific Confluent Cloud environment IDs to filter. When provided, only these environments will be used to fetch Kafka cluster credentials. Use commas to separate multiple environment IDs. Leave blank or unset to use all available environments.	env-123,env-456	Empty (all environments)	No
PRINCIPAL_ID	String	Confluent Cloud principal ID (user or service account) for API key creation.	u-abc123 or sa-xyz789	None	Yes
KAFKA_CLUSTER_FILTER	Comma-separated String	A list of specific Kafka cluster IDs to filter. When provided, only these Kafka clusters will be analyzed. Use commas to separate multiple cluster IDs. Leave blank or unset to analyze all available clusters.	lkc-123,lkc-456	Empty (all clusters)	No
CONFLUENT_CLOUD_CREDENTIAL	JSON Object	Contains authentication credentials for Confluent Cloud API access. Must include <code>confluent_cloud_api_key</code> and <code>confluent_cloud_api_secret</code> fields for authenticating with Confluent Cloud services.	{"confluent_cloud_api_key": "CKABCD123456", "confluent_cloud_api_secret": "xyz789secretkey"}	None	Yes (if not in Manager)
USE_AWS_SECRETS_MANAGER	Boolean	Controls whether to retrieve credentials from AWS Secrets Manager instead of using direct environment variables. When <code>True</code> , credentials are fetched from AWS Secrets Manager using the paths specified in other variables.	True or False	False	No

Environment Variable Name	Type	Description	Example	Default	Required
CONFLUENT_CLOUD_API_SECRET_PATH	JSON Object	AWS Secrets Manager configuration for Confluent Cloud credentials. Contains <b>region_name</b> (AWS region) and <b>secret_name</b> (name of the secret in AWS Secrets Manager). Only used when <b>USE_AWS_SECRETS_MANAGER</b> is <b>True</b> .	<code>{"region_name": "us-east-1", "secret_name": "confluent-cloud-api-credentials"}</code>	None	Yes (if <b>USE_AWS_</b> is <b>True</b> )

1.2.3 Using the AWS Secrets Manager (optional)

If you use **AWS Secrets Manager** to manage your secrets, set the **USE\_AWS\_SECRETS\_MANAGER** variable to **True** and the tool will retrieve the secrets from AWS Secrets Manager using the names provided in **CONFLUENT\_CLOUD\_API\_KEY\_AWS\_SECRETS**.

The code expects the **CONFLUENT\_CLOUD\_API\_KEY\_AWS\_SECRETS** to be stored in JSON format with these keys:

- `confluent_cloud_api_key`
- `confluent_cloud_api_secret`

1.3 Run the Tool

Navigate to the Project Root Directory

Open your Terminal and navigate to the root folder of the `kafka_cluster-topic-key_distribution_analyzer-tool/` repository that you have cloned. You can do this by executing:

```
cd path/to/kafka_cluster-topic-key_distribution_analyzer-tool/
```

Replace `path/to/` with the actual path where your repository is located.

Then enter the following command below to run the tool:

```
uv run streamlit run src/tool.py
```

1.3.1 Did you notice we prefix `uv run` to `streamlit run src/tool.py`?

You maybe asking yourself why. Well, `uv` is an incredibly fast Python package installer and dependency resolver, written in **Rust**, and designed to seamlessly replace `pip`, `pipx`, `poetry`, `pyenv`, `twine`, `virtualenv`, and more in your workflows. By prefixing `uv run` to a command, you're ensuring that the command runs in an optimal Python environment.

Now, let's go a little deeper into the magic behind `uv run`:

- When you use it with a file ending in `.py` or an HTTP(S) URL, `uv` treats it as a script and runs it with a Python interpreter. In other words, `uv run file.py` is equivalent to `uv run python file.py`. If you're working with a URL, `uv` even downloads it temporarily to execute it. Any inline dependency metadata is installed into an isolated, temporary environment—meaning zero leftover mess! When used with `-`, the input will be read from `stdin`, and treated as a Python script.
- If used in a project directory, `uv` will automatically create or update the project environment before running the command.
- Outside of a project, if there's a virtual environment present in your current directory (or any parent directory), `uv` runs the command in that environment. If no environment is found, it uses the interpreter's environment.

So what does this mean when we put `uv run` before `streamlit run src/tool.py`? It means `uv` takes care of all the setup—fast and seamless—right in your local environment. If you think AI/ML is magic, the work the folks at **Astral** have done with `uv` is pure wizardry!

Curious to learn more about **Astral's** `uv`? Check these out:

- Documentation: Learn about `uv`.
- Video: [uv IS THE FUTURE OF PYTHON PACKING!](#)

If you have Kafka connectivity issues, you can verify connectivity using the following command:

1.3.2 Troubleshoot Connectivity Issues (if any)

To verify connectivity to your Kafka cluster, you can use the `kafka-topics.sh` command-line tool. First, download the Kafka binaries from the **Apache Kafka website** and extract them. Navigate to the `bin` directory of the extracted Kafka folder. Second, create a `client.properties` file with your Kafka credentials:

```
# For SASL_SSL (most common for cloud services)
security.protocol=SASL_SSL
sasl.mechanism=PLAIN
sasl.jaas.config=org.apache.kafka.common.security.plain.PlainLoginModule required \
```

```
username="<YOUR_KAFKA_API_KEY>" \
password="<YOUR_KAFKA_API_SECRET>";

# Additional SSL settings if needed
ssl.endpoint.identification.algorithm=https
```

Finally, run the following command to list all topics in your Kafka cluster:


```
./kafka-topics.sh --list --bootstrap-server <YOUR_BOOTSTRAP_SERVER_URI> --command-config ./client.properties
```

If the connection is successful, you should see a list of topics in your Kafka cluster. If you encounter any errors, double-check your credentials and network connectivity.

1.4 The Results

Once the tool completes its analysis, it will display a dashboard with visualizations and metrics for each partitioning strategy tested. You will see bar charts showing the distribution of messages across partitions for each strategy, along with summary statistics like standard deviation and coefficient of variation.

1.4.1 Example of Hot Key Data Skew Simulation Results



1.4.2 Example of Normal Key Distribution Simulation Results

Key Distribution Analyzer Tool Dashboard

This teaching tool shows you how the different key patterns, key simulation strategies, and partition strategies affect the key distribution.

Choose the Environment:  
dev

Choose the dev's Kafka Cluster:  
kafka-broker-dev-us-east-1

Enter your Kafka Producer topic name:  
\_j3\_key\_distribution\_9

Key Pattern:  
[tenant\_id:'object\_id']

Key Simulation:  
Normal


Partition Count:  
10

Record Count:  
1,000

Run Key Distribution Analysis Tests

Analysis tests are complete

Kafka Partition Strategy Comparison using a Normal Key Simulation



MurmurHash2 is a non-cryptographic hash function that was created by Austin Appleby in 2006, produces 32-bit hash values, is extremely fast (3-5x faster than MD5), has excellent distribution properties, and is used by Kafka, Redis, Cassandra, and many others. For more information, see the [MurmurHash Wikipedia page](#).

Round Robin is the simplest partitioning strategy that ignores the message key completely, distributes messages sequentially across partitions, and cycles through partitions in order: 0 → 1 → 2 → 3 → ... → 0 (repeats). The name Round Robin comes from a 16th-century French term meaning 'ribbon round' - signing documents in a circle so no one appears first!

Sticky partitioning is a strategy that assigns messages to a single partition for a batch, sticks to that partition until the batch is full or a timeout occurs, and then switches to a new partition for the next batch. This approach reduces the overhead of frequent partition switching and improves throughput while still providing some level of distribution across partitions.

Range-Based Custom partitioning is a strategy that assigns messages to partitions based on predefined key ranges, sorts unique keys and divides them into ranges corresponding to each partition, and ensures that similar keys are grouped together in the same partition. This approach is useful for scenarios where key locality is important, such as time-series data or ordered processing.

Custom partitioning is a simple strategy that uses Python's built-in hash function to compute a hash value for each key, applies a modulo operation with the number of partitions to determine the target partition, and distributes messages based on the computed partition. This approach is straightforward but may not provide optimal distribution compared to more sophisticated hashing algorithms.

Partition Strategy Metrics Summary

Partition Strategy	Total Records	Average per Partition	Standard Deviation	Coefficient of Variation (%)	Quality
Default (Murmur Hash 2)	1000	100.0	27.89	27.9	Needs Improvement
Round Robin	1000	100.0	0.00	0.0	Good
Sticky	1000	100.0	94.16	94.2	Needs Improvement
Range Based Custom	1000	100.0	31.62	31.6	Needs Improvement
Custom	1000	100.0	37.71	37.7	Needs Improvement

Standard Deviation (SD) measures the amount of variation or dispersion in a set of values. A low SD indicates that the values tend to be close to the mean, while a high SD indicates that the values are spread out over a wider range.

Coefficient of Variation (CV) is a standardized measure of dispersion of a probability distribution or frequency distribution. It is often expressed as a percentage and is defined as the ratio of the standard deviation to the mean. A lower CV indicates a more uniform distribution, while a higher CV indicates greater variability.

In general, a CV less than 20% is considered good, indicating a relatively uniform distribution across partitions. A CV greater than 20% suggests that the distribution may be uneven and could benefit from optimization.

Quality indicators: Good (CV < 20%), Needs Improvement (CV > 20%)

Note: These metrics are based on the produced records and may vary with different key patterns, record counts, and partition counts. They provide insights into how well each partitioning strategy distributes messages across partitions.

Cleanup Resources before Closing the Tool

1.4.3 Example of Low Cardinality Key Distribution Simulation Results

Key Distribution Analyzer Tool Dashboard

This teaching tool shows you how the different key patterns, key simulation strategies, and partition strategies affect the key distribution.

Choose the Environment:

dev

Choose the dev's Kafka Cluster:

kafka-broker-dev-us-east-1

Enter your Kafka Producer topic name:

\_j3\_key\_distribution\_9

Key Pattern:

["tenant\_id","object\_id"]

Key Simulation:

Less Repetition

Partition Count:

10


Record Count:

1,000

Run Key Distribution Analysis Tests

Analysis tests are complete

Kafka Partition Strategy Comparison using a Less Repetition Key Simulation



MurmurHash2 is a non-cryptographic hash function that was created by Austin Appleby in 2008, produces 32 bit hash values, is extremely fast (3-5x faster than MD5), has excellent distribution properties, and is used by Kafka, Redis, Cassandra, and many others. For more information, see the [MurmurHash Wikipedia page](#).

Round Robin is the simplest partitioning strategy that ignores the message key completely, distributes messages sequentially across partitions, and cycles through partitions in order: 0 + 1 + 2 + 3 + ... + 9 (repeats). The name Round Robin comes from a 16th-century French term meaning 'ribbon round' - signing documents in a circle so no one appears first!

Sticky partitioning is a strategy that assigns messages to a single partition for a batch, sticks to that partition until the batch is full or a timeout occurs, and then switches to a new partition for the next batch. This approach reduces the overhead of frequent partition switching and improves throughput while still providing some level of distribution across partitions.

Range-Based Custom partitioning is a strategy that assigns messages to partitions based on predefined key ranges, sorts unique keys and divides them into ranges corresponding to each partition, and ensures that similar keys are grouped together in the same partition. This approach is useful for scenarios where key locality is important, such as time-series data or ordered processing.

Custom partitioning is a simple strategy that uses Python's built-in hash function to compute a hash value for each key, applies a modulo operation with the number of partitions to determine the target partition, and distributes messages based on the computed partition. This approach is straightforward but may not provide optimal distribution compared to more sophisticated hashing algorithms.

Partition Strategy Metrics Summary

Partition Strategy	Total Records	Average per Partition	Standard Deviation	Coefficient of Variation (%)	Quality
Default (Murmur Hash 2)	1000	100.0	11.60	11.6	Good
Round Robin	1000	100.0	0.00	0.0	Good
Sticky	1000	100.0	66.50	66.5	Needs Improvement
Range Based Custom	1000	100.0	3.36	3.2	Good
Custom	1000	100.0	10.51	10.5	Good

Standard Deviation (SD) measures the amount of variation or dispersion in a set of values. A low SD indicates that the values tend to be close to the mean, while a high SD indicates that the values are spread out over a wider range.

Coefficient of Variation (CV) is a standardized measure of dispersion of a probability distribution or frequency distribution. It is often expressed as a percentage and is defined as the ratio of the standard deviation to the mean. A lower CV indicates a more uniform distribution, while a higher CV indicates greater variability.

In general, a CV less than 20% is considered good, indicating a relatively uniform distribution across partitions. A CV greater than 20% suggests that the distribution may be uneven and could benefit from optimization.

Quality indicators: Good (CV < 20%), Needs Improvement (CV > 20%)

Note: These metrics are based on the produced records and may vary with different key patterns, record counts, and partition counts. They provide insights into how well each partitioning strategy distributes messages across partitions.

Cleanup Resources before Closing the Tool

1.4.4 Example of High Cardinality Key Distribution Simulation Results

Key Distribution Analyzer Tool Dashboard

This teaching tool shows you how the different key patterns, key simulation strategies, and partition strategies affect the key distribution.

Choose the Environment:

dev

Choose the dev's Kafka Cluster:

kafka-broker-dev-us-east-1

Enter your Kafka Producer topic name:

\_j3\_key\_distribution\_9

Key Pattern:

['tenant\_id','object\_id']

Key Simulation:

More Repetition

Partition Count:

10

Record Count:

1,000

Run Key Distribution Analysis Tests

Analysis tests are complete

Kafka Partition Strategy Comparison using a More Repetition Key Simulation

The figure displays six bar charts comparing different Kafka partitioning strategies. Each chart shows the number of messages per partition (y-axis) against the partition index (x-axis). A red dashed line indicates the average message count per partition. A green icon and CV (Coefficient of Variation) percentage are shown for each strategy.

- Default (Murmur Hash 2):** CV: 105.4% (Needs Improvement). Messages are unevenly distributed across 10 partitions.
- Round Robin:** CV: 0.0% (Good). Messages are perfectly evenly distributed across 10 partitions.
- Sticky:** CV: 94.0% (Needs Improvement). Messages are unevenly distributed across 10 partitions.
- Range Based Custom:** CV: 105.4% (Needs Improvement). Messages are unevenly distributed across 10 partitions.
- Partition Custom:** CV: 81.6% (Needs Improvement). Messages are unevenly distributed across 10 partitions.
- Custom:** CV: 105.4% (Needs Improvement). Messages are unevenly distributed across 10 partitions.

**MurmurHash2** is a non-cryptographic hash function that was created by Austin Appleby in 2008, produces 32-bit hash values, is extremely fast (3-5x faster than MD5), has excellent distribution properties, and is used by Kafka, Redis, Cassandra, and many others. For more information, see the [MurmurHash Wikipedia page](#).

**Round Robin** is the simplest partitioning strategy that ignores the message key completely, distributes messages sequentially across partitions, and cycles through partitions in order: 0 + 1 + 2 + 3 + ... + 0 (repeats). The name **Round Robin** comes from a 16th-century French term meaning 'ribbon round' - signing documents in a circle so no one appears first!

**Sticky** partitioning is a strategy that assigns messages to a single partition for a batch, sticks to that partition until the batch is full or a timeout occurs, and then switches to a new partition for the next batch. This approach reduces the overhead of frequent partition switching and improves throughput while still providing some level of distribution across partitions.

**Range-Based Custom** partitioning is a strategy that assigns messages to partitions based on predefined key ranges, sorts unique keys and divides them into ranges corresponding to each partition, and ensures that similar keys are grouped together in the same partition. This approach is useful for scenarios where key locality is important, such as time-series data or ordered processing.

**Custom** partitioning is a simple strategy that uses Python's built-in hash function to compute a hash value for each key, applies a modulo operation with the number of partitions to determine the target partition, and distributes messages based on the computed partition. This approach is straightforward but may not provide optimal distribution compared to more sophisticated hashing algorithms.

Partition Strategy Metrics Summary

Partition Strategy	Total Records	Average per Partition	Standard Deviation	Coefficient of Variation (%)	Quality
Default (Murmur Hash 2)	1000	100.0	105.41	105.4	Needs Improvement
Round Robin	1000	100.0	0.00	0.0	Good
Sticky	1000	100.0	94.05	94.0	Needs Improvement
Range Based Custom	1000	100.0	105.41	105.4	Needs Improvement
Custom	1000	100.0	81.65	81.6	Needs Improvement

**Standard Deviation (SD)** measures the amount of variation or dispersion in a set of values. A low SD indicates that the values tend to be close to the mean, while a high SD indicates that the values are spread out over a wider range.

**Coefficient of Variation (CV)** is a standardized measure of dispersion of a probability distribution or frequency distribution. It is often expressed as a percentage and is defined as the ratio of the standard deviation to the mean. A lower CV indicates a more uniform distribution, while a higher CV indicates greater variability.

In general, a CV less than 20% is considered good, indicating a relatively uniform distribution across partitions. A CV greater than 20% suggests that the distribution may be uneven and could benefit from optimization.

**Quality indicators:** Good (CV < 20%), Needs Improvement (CV ≥ 20%)

*Note: These metrics are based on the produced records and may vary with different key patterns, record counts, and partition counts. They provide insights into how well each partitioning strategy distributes messages across partitions.*

Cleanup Resources before Closing the Tool

1.4.5 Example of No Repetition Key Distribution Simulation Results

7 / 10

Key Distribution Analyzer Tool Dashboard

This teaching tool shows you how the different key patterns, key simulation strategies, and partition strategies affect the key distribution.

Choose the Environment:

dev

Choose the dev's Kafka Cluster:

kafka-broker-dev-us-east-1

Enter your Kafka Producer topic name:

\_j3\_key\_distribution\_9

Key Pattern:

["tenant\_id","object\_id"]

Key Simulation:

No Repetition

Partition Count:

10


Record Count:

1,000

Run Key Distribution Analysis Tests

Analysis tests are complete

Kafka Partition Strategy Comparison using a No Repetition Key Simulation



**MurmurHash2** is a non-cryptographic hash function that was created by [Austin Appleby](#) in 2008, produces 32-bit hash values, is extremely fast (3-5x faster than MD5), has excellent distribution properties, and is used by *Kafka*, *Redis*, *Cassandra*, and many others. For more information, see the [MurmurHash Wikipedia page](#).

**Round Robin** is the simplest partitioning strategy that ignores the message key completely, distributes messages sequentially across partitions, and cycles through partitions in order: 0 + 1 + 2 + 3 + ... + 0 (repeats). The name **Round Robin** comes from a 16th-century French term meaning "ribbon round" - signing documents in a circle so no one appears first!

**Sticky** partitioning is a strategy that assigns messages to a single partition for a batch, sticks to that partition until the batch is full or a timeout occurs, and then switches to a new partition for the next batch. This approach reduces the overhead of frequent partition switching and improves throughput while still providing some level of distribution across partitions.

**Range-Based Custom** partitioning is a strategy that assigns messages to partitions based on predefined key ranges, sorts unique keys and divides them into ranges corresponding to each partition, and ensures that similar keys are grouped together in the same partition. This approach is useful for scenarios where key locality is important, such as time-series data or ordered processing.

**Custom** partitioning is a simple strategy that uses Python's built-in hash function to compute a hash value for each key, applies a modulo operation with the number of partitions to determine the target partition, and distributes messages based on the computed partition. This approach is straightforward but may not provide optimal distribution compared to more sophisticated hashing algorithms.

Partition Strategy	Total Records	Average per Partition	Standard Deviation	Coefficient of Variation (%)	Quality
Default (Murmur Hash 2)	1000	100.0	11.60	11.6	Good
Round Robin	1000	100.0	0.00	0.0	Good
Sticky	1000	100.0	66.67	66.7	Needs Improvement
Range Based Custom	1000	100.0	3.16	3.2	Good
Custom	1000	100.0	10.51	10.5	Good

**Standard Deviation (SD)** measures the amount of variation or dispersion in a set of values. A low SD indicates that the values tend to be close to the mean, while a high SD indicates that the values are spread out over a wider range.

**Coefficient of Variation (CV)** is a standardized measure of dispersion of a probability distribution or frequency distribution. It is often expressed as a percentage and is defined as the ratio of the standard deviation to the mean. A lower CV indicates a more uniform distribution, while a higher CV indicates greater variability.

In general, a CV less than 20% is considered good, indicating a relatively uniform distribution across partitions. A CV greater than 20% suggests that the distribution may be uneven and could benefit from optimization.

**Quality indicators:** Good (CV < 20%), Needs Improvement (CV ≥ 20%)

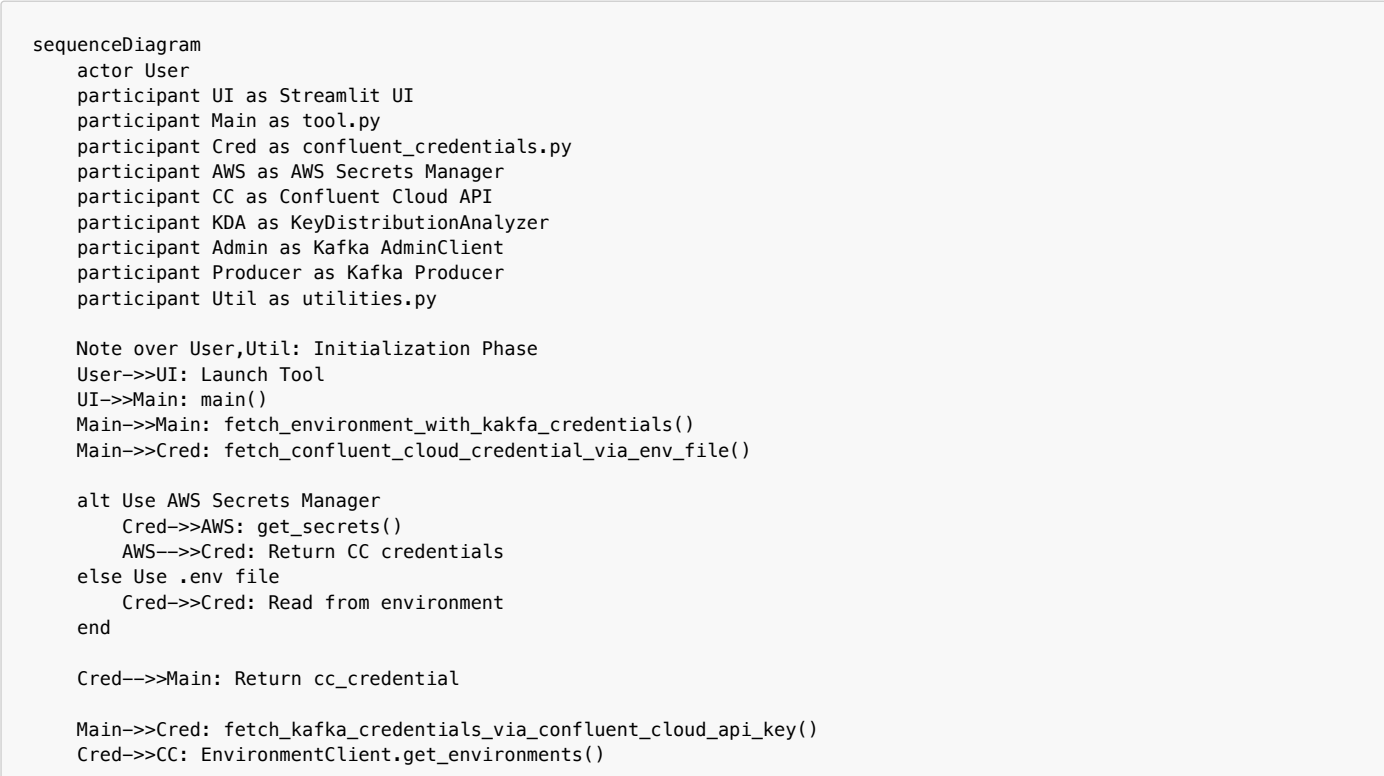
*Note: These metrics are based on the produced records and may vary with different key patterns, record counts, and partition counts. They provide insights into how well each partitioning strategy distributes messages across partitions.*

Cleanup Resources before Closing the Tool

2.0 How the Tool Works

2.1 End-to-End Flow

The following sequence diagram illustrates the interactions between the user, Streamlit UI, and various components of the tool during its execution:





```

CC-->>Cred: Return environments

loop For each environment
  Cred->>CC: EnvironmentClient.get_kafka_clusters()
  CC-->>Cred: Return kafka_clusters

  loop For each Kafka cluster
    Cred->>CC: IamClient.create_api_key()
    CC-->>Cred: Return API key pair
    Cred->>Cred: Store kafka_credentials
  end
end

Cred-->>Main: Return environments, kafka_clusters, kafka_credentials
Main-->>UI: Display environment & cluster selection

Note over User,Util: Configuration Phase
User->>UI: Select environment
User->>UI: Select Kafka cluster
User->>UI: Enter topic name
User->>UI: Configure key pattern
User->>UI: Select key simulation type
User->>UI: Set partition count
User->>UI: Set record count
User->>UI: Click "Run Key Distribution Analysis Tests"

Note over User,Util: Execution Phase
UI->>Main: run_tests()
Main->>KDA: Initialize KeyDistributionAnalyzer
KDA->>KDA: Setup AdminClient config
KDA->>KDA: Setup Producer config
KDA->>KDA: Setup Consumer config

Main->>KDA: run_test()
KDA->>UI: progress_bar.progress(0.125)

KDA->>Util: create_topic_if_not_exists()
Util->>Admin: list_topics()
Admin-->>Util: Return topic list

alt Topic exists
  Util->>Admin: delete_topics()
  Admin-->>Util: Confirm deletion
end

Util->>Admin: create_topics()
Admin-->>Util: Confirm creation
Util-->>KDA: Return success

KDA->>UI: progress_bar.progress(0.25)
KDA->>KDA: __produce_test_records()

loop For each record
  KDA->>KDA: Generate key based on simulation type
  alt Normal
    KDA->>KDA: key = pattern + (id % 100)
  else Less Repetition
    KDA->>KDA: key = pattern + (id % 1000)
  else More Repetition
    KDA->>KDA: key = pattern + (id % 10)
  else No Repetition
    KDA->>KDA: key = pattern + id
  else Hot Key Data Skew
    alt 80% of records
      KDA->>KDA: key = "hot-key"
    else 20% of records
      KDA->>KDA: key = "cold-key-" + id
    end
  end

  KDA->>Producer: produce(topic, key, value)
  Producer->>KDA: __delivery_callback()
  KDA->>KDA: Store key in partition_mapping
end

KDA->>Producer: flush()

KDA->>UI: progress_bar.progress(0.375)
KDA->>KDA: __analyze_distribution()
KDA->>KDA: Calculate partition record counts
KDA->>KDA: Calculate key pattern distribution

```

```

KDA-->>UI: progress_bar.progress(0.5)
KDA-->>KDA: __test_partition_strategies()

par Test All Strategies
  KDA-->>KDA: __murmur2_hash_strategy()
  KDA-->>KDA: __round_robin_strategy()
  KDA-->>KDA: __sticky_strategy()
  KDA-->>KDA: __range_based_customer_strategy()
  KDA-->>KDA: __custom_strategy()
end

KDA-->>UI: __visualize_strategy_comparison()
UI-->>UI: Display Plotly charts
UI-->>UI: Display metrics summary

KDA-->>UI: progress_bar.progress(0.625)
KDA-->>KDA: __test_hash_distribution()
KDA-->>KDA: Calculate theoretical distribution

KDA-->>UI: progress_bar.progress(0.75)
KDA-->>KDA: Compare actual vs theoretical

KDA-->>UI: progress_bar.progress(0.875)
KDA-->>KDA: Calculate quality metrics
KDA-->>KDA: Compute mean, std dev, CV

KDA-->>UI: progress_bar.progress(1.0)
KDA-->>Main: Return distribution_results
Main-->>>UI: Display success & balloons

Note over User,Util: Cleanup Phase
User-->>UI: Click "Cleanup Resources"
UI-->>Main: delete_all_kafka_credentials_created()

loop For each kafka_credential
  Main-->>CC: IAMClient.delete_api_key()
  CC-->>>Main: Confirm deletion
end

Main-->>>UI: Display success message
UI-->>>User: Tool ready to close

```

### 3.0 Resources

- [The Importance of Standardized Hashing Across Producers](#)
- [librdkafka Configuration](#)